

Secure Computation: Chaining Garbled Circuits:

Alex Ledger

Reed College

April 12, 2016

Overview

1. Secure Computation
2. Garbled Circuits
3. Chaining Garbled Circuits
4. My work: SCMC and Programming

Goals and Notes

1. Understand the high level idea of secure computation
 - ▶ Maybe you'll run into a situation where it's useful, and it can help solve an otherwise unsolvable, real world problem.
2. Understand garbled circuits - the most basic construction
3. Understand chaining - how we (my thesis) made secure computation faster

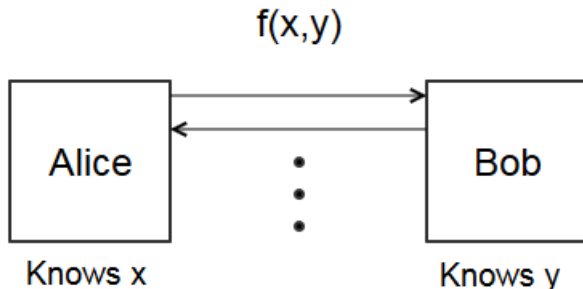
Before we start:

- ▶ I will not emphasize security
- ▶ Lots of notation: ask me if I brush over something, or you forget what something means
- ▶ Lots of moving parts

The Millionaire Problem

- ▶ Alice and Bob want to determine who is wealthier.
- ▶ They do want to disclose their wealth to each other.
 - ▶ Alice has \$ x , Bob has \$ y .
 - ▶ Alice should not learn anything about y .
 - ▶ Bob should not learn anything about x .


$$f(x, y) = \begin{cases} \text{Alice}, & y \leq x; \\ \text{Bob}, & y > x. \end{cases} \quad (1)$$



Security Properties

- ▶ Privacy of inputs
 - ▶ Alice and Bob do not learn anything about the other's input.
 - ▶ Except for info that is inferable from x and $f(x, y)$.
 - ▶ Bob should not learn that $1,000,000 < x \leq 2,000,000$.
 - ▶ But if $y < x$ and $y = 2,000$, then he learns $x < 2,000$.
- ▶ Correctness
 - ▶ Alice and Bob receive $f(x, y)$.
 - ▶ As opposed to some value near $f(x, y)$.
 - ▶ Or not receiving a value at all
- ▶ Semi-honest
 - ▶ We assume that each party obeys the protocol, but attempts to learn extra information from its interactions


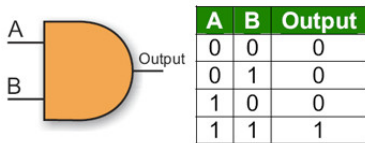
Security



A secure computation protocol is secure if Alice and Bob learn the same information in the real world as they would in ideal world.

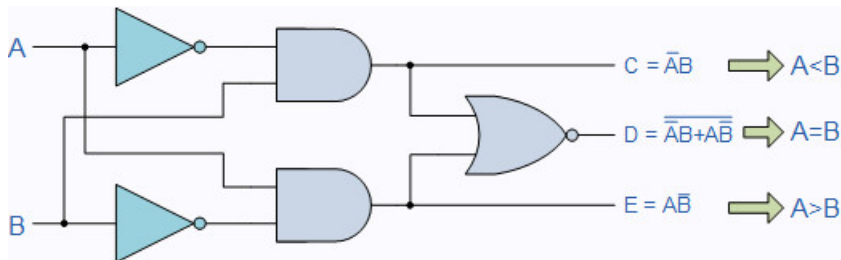
Boolean Circuits

- ▶ We encode a function f into a circuit C .
- ▶ Circuit C is made of AND, XOR and NOT gates.
- ▶ Each gate has two input wires and a single output wire




Boolean Circuits

- ▶ Any function can be encoded into a circuit.
- ▶ Here is the less than circuit.




Encryption

Encryption




Decryption



Oblivious Transfer (OT) in brief

- ▶ Alice potentially sends either m_0 or m_1 to Bob.
- ▶ Bob, without seeing m_0 or m_1 , decides that he wants m_b .
- ▶ Bob receives m_b .
- ▶ Property 1: Alice does not know which message Bob received.
- ▶ Property 2: Bob doesn't anything about m_{1-b} .



Hash Function in brief

- ▶ Hash function $H : \{0, 1\}^* \rightarrow \{0, 1\}^{128}$
- ▶ For our purposes, H maps any string to a uniform, random 128-bit string.
- ▶ A.k.a. H is a random oracle.

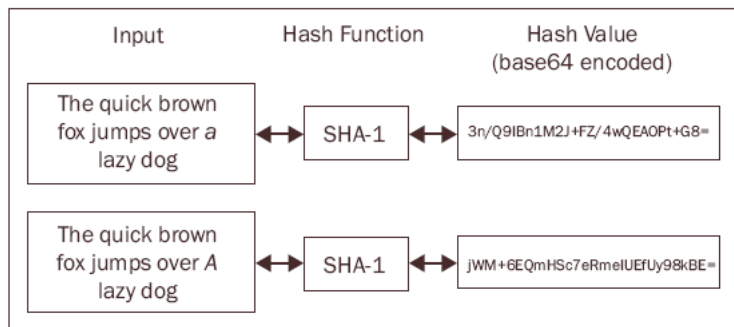



Figure 17: Hash Function

Roadmap of Garbled Circuits


1. Alice (garbler) *garbles* the AND gate.
2. Alice sends the *garbled table* of the gate and some keys to Bob.
3. Bob (evaluator) *evaluates* the gate.



Garbling a gate 1

Step 1. Alice assigns *wire labels* to each wire.


- ▶ For each wire in the circuit, assign two random labels to each wire
- ▶ Wire A has two wire labels A_0 and A_1 .
- ▶ We say A_0 *semantically represents* 0, and A_1 *semantically represents* 1.
- ▶ And A_0 and A_1 are sampled uniform randomly from $\{0, 1\}^n$.
 - ▶ Generally use AES-128 for encryption, so $n = 128$.



Garbling a gate 2

Step 2. Alice constructs garbled table.

- ▶ Encrypt wire labels for C , C_0 and C_1 , using wire labels of A and B .
- ▶ Randomly permute table



A	B	Encryption
A_0	B_0	$\text{Enc}_{A_0, B_0}(C_0)$
A_1	B_0	$\text{Enc}_{A_1, B_0}(C_0)$
A_0	B_1	$\text{Enc}_{A_0, B_1}(C_0)$
A_1	B_1	$\text{Enc}_{A_1, B_1}(C_1)$

Garbling a gate 3

Step 3. Send garbled table to Bob.

$\text{Enc}_{A_0, B_0}(C_0)$
$\text{Enc}_{A_1, B_0}(C_0)$
$\text{Enc}_{A_0, B_1}(C_0)$
$\text{Enc}_{A_1, B_1}(C_1)$

Garbling a gate 4

Step 4. Alice sends input wire labels to Bob.

- ▶ Suppose $x \in \{0, 1\}$ is Alice's input, and $y \in \{0, 1\}$ is Bob's input.
- ▶ Alice sends A_x to Bob.
- ▶ Alice sends B_y to Bob via Oblivious Transfer
 - ▶ The wire labels corresponding to her inputs.
- ▶ Bob has:

Garbled Table
$\text{Enc}_{A_0, B_0}(C_0)$
$\text{Enc}_{A_1, B_0}(C_0)$
$\text{Enc}_{A_0, B_1}(C_0)$
$\text{Enc}_{A_1, B_1}(C_1)$

Input Labels
A_*
B_y

Garbling a gate 5

Step 5. Bob evaluates the circuit

- ▶ Bob has the garbled table, A_x and B_y .
- ▶ Bob trial decrypts each row of the garbled table, until an encryption succeeds.
- ▶ Bob acquires $C_{x \wedge y}$.
- ▶ Bob has:

Garbled Table
$\text{Enc}_{A_0, B_0}(C_0)$
$\text{Enc}_{A_1, B_0}(C_0)$
$\text{Enc}_{A_0, B_1}(C_0)$
$\text{Enc}_{A_1, B_1}(C_1)$

Input Labels
A_*
B_y

Output Label
$C_{x \wedge y}$

Garbling a gate 6

Step 6. Bob gets a final answer.

- ▶ Alice sends $\text{Enc}_{C_0}(0)$ and $\text{Enc}_{C_1}(1)$ to Bob.
- ▶ Bob trial decrypts these with $C_{x \wedge y}$.
- ▶ One will succeed, and Bob will acquire $z = x \wedge y$.
- ▶ So Bob knows z , but not x !

Garbled Table
$\text{Enc}_{A_0, B_0}(C_0)$
$\text{Enc}_{A_1, B_0}(C_0)$
$\text{Enc}_{A_0, B_1}(C_0)$
$\text{Enc}_{A_1, B_1}(C_1)$


Input Labels
A_*
B_y

Output Label
$C_{x \wedge y}$

Output Map
$\text{Enc}_{C_0}(0)$
$\text{Enc}_{C_1}(1)$


Security Considerations

- ▶ Think about what Alice acquires:
 - ▶ Well Bob had no input in this case, so ...
 - ▶ In general, Alice generates objects and sends them to Bob
 - ▶ So she doesn't have many opportunities to learn about Bob's input.
- ▶ Think about what Bob acquires:
 1. Garbled table
 2. Input wire labels: A_a and B_b
 3. Encryptions of output: $\text{Enc}_{C_0}(0)$ and $\text{Enc}_{C_1}(1)$
- ▶ Can Bob learn anything about a or b ?
 - ▶ If he could decrypt another row of the garbled table, then we would learn something.
 - ▶ But he can't, because he doesn't have the keys.
 - ▶ Not really, because everything is encrypted.




Extending a garbled gate into a garbled circuit

- ▶ To operate on a more complex function, the operation is recursed.
- ▶ The output wires of the first gates are used as inputs to subsequent gates.
- ▶ Alice only sends output maps for the final gates.



The Garbled Circuit Protocol



The cost of garbled circuits

- ▶ Alice sends 4 ciphertexts per gate, since the garbled table has 4 rows, to Bob.
- ▶ This is the bandwidth, or the amount of the communication required.
- ▶ Based on empirical work, bandwidth is the biggest bottleneck in garbled circuits.
- ▶ So reducing the size of the garbled table is of utmost priority.

Reducing the size of the garbled table with Free XOR

- ▶ Let $\Delta \leftarrow \{0, 1\}^n$ be fixed globally in a circuit.
- ▶ For each input wire A , sample a single ciphertext; call it A .
 - ▶ The *zeroth* wire label is A .
 - ▶ The *first* wire label is $A \oplus \Delta$.
- ▶ Set output wire C of a gate to be $A \oplus B$.
 - ▶ The xor of its input wires.
- ▶ Bob *evaluates* the XOR gate by XORing the two input labels.

$$(A \oplus a\Delta) \oplus (B \oplus b\Delta) = A \oplus B \oplus (a \oplus b)\Delta$$

- ▶ So XOR gates do not require a garbled table, aka they're Free.

Offline/Online

- ▶ Imagine two banks use secure computation during their daily operations
 - ▶ E.g. ATM operations
- ▶ At night - the offline phase - they exchange garbled circuits (the garbled tables)
 - ▶ So garbled circuits are ATM transactions
- ▶ During the day - the online phase - they exchange input wire labels and use the pre-exchanged garbled circuits.
- ▶ The computation is fast!
- ▶ Problems
 - ▶ Functions are decided at night - no room for flexibility
 - ▶ Input size is fixed at night


ATM Transaction

Insurance Operation

Statistical Operation


Chaining Garbled Circuits

- ▶ Goal: improve offline/online garbled circuits by adding flexibility
- ▶ Key observation: many useful functions in the real world are composed of small, standard components.
 - ▶ E.g. addition, subtraction, matrix operations
 - ▶ Leveshtein distance algorithm - a dynamic algorithm
 - ▶ Encryption is a common component (banks)
- ▶ Idea: chain garbled circuits together
 - ▶ Take the output of one garbled circuit and plug it into another garbled circuit




How to chain garbled circuits

- ▶ Suppose that we are chaining a garbled circuit with output wire A to garbled circuit with input wire X .
- ▶ We want $A \rightarrow X$ and $X \oplus \Delta \rightarrow X \oplus \Delta$.
- ▶ Straightforward:
 - ▶ Alice sends Bob $L_{AX} = A \oplus X$
 - ▶ Bob sets $X_* \leftarrow A_* \oplus L_{AX}$



How to chain garbled circuits

- ▶ Suppose that we are chaining a garbled circuit with output wire A to garbled circuit with input wire X .
- ▶ We want $A \rightarrow X$ and $X \oplus \Delta \rightarrow X \oplus \Delta$.
- ▶ Straightforward:
 - ▶ Alice sends Bob $L_{AX} = A \oplus X$
 - ▶ Bob sets $X_* \leftarrow A_* \oplus L_{AX}$



- ▶ Is this secure?

Efficiency of chaining garbled circuits

- ▶ In online phase, chaining requires communicating a ciphertext per chain
- ▶ This can be a lot of communication
 - ▶ Suppose 100 by 100 matrix with entries of max value 2^8
 - ▶ Then requires 8,000 ciphertexts
 - ▶ That's 256 kBs.
- ▶ I came up with a method to reduce this to a single ciphertext per data object


Single Communication Multiple Connections SCMC

- ▶ Problem: number of ciphertexts communicated scales linearly with the number of chains
- ▶ Key observation:
 - ▶ the chaining is predictable, consecutive wires are chained to consecutive wires
 - ▶ Pieces of data move around in a function, like matrices or numbers or text
- ▶ Solution: give consecutive wire labels a predictable pattern



Single Communication Multiple Connections SCMC

- ▶ For each piece of data, sample a label A .
- ▶ Set i th wire label of data to $A \oplus H(i)$.
- ▶ Where H is a hash function



- ▶ Technically, we set the links to $A \oplus H(T \oplus (i||b))$

My thesis work

- ▶ I implemented chaining and SCMC in a program called CompGC
- ▶ Written in C
- ▶ It does the entire garbled circuits protocol
- ▶ Takes circuits and inputs and securely computed the output
- ▶ It's fast
 - ▶ Leventshein-60 with SCMC: ≈ 750 ms
 - ▶ Without chaining: ≈ 10 s

Conclusion

We talked about:

- ▶ Secure computation
- ▶ Garbled circuits
- ▶ Chaining garbled circuits
- ▶ SCMC
- ▶ CompGC

Thank you:

- ▶ Adam Groce
- ▶ Alex Malozemoff and Arkady Yerukhimovich
- ▶ Catdog and friends